

EECE5640 Homework 4

Question 1

a.)

Please see the details in the code part -> q1/q1a/q1a.c

The result is as follow:

```
[[li.junce@login-00 q1a]$ ls
q1a q1a.c q1a.out q1a.script
[[li.junce@login-00 q1a]$ more q1a.out
The Node c0376's Process 0 increases the count from 0 to 1.
The Node c0376's Process 1 increases the count from 1 to 2.
The Node c0376's Process 2 increases the count from 2 to 3.
The Node c0376's Process 3 increases the count from 3 to 4.
The Node c0376's Process 4 increases the count from 4 to 5.
The Node c0376's Process 5 increases the count from 5 to 6.
The Node c0376's Process 6 increases the count from 6 to 7.
The Node c0376's Process 7 increases the count from 7 to 8.
The Node c0376's Process 8 increases the count from 8 to 9.
The Node c0376's Process 9 increases the count from 9 to 10.
The Node c0376's Process 10 increases the count from 10 to 11.
The Node c0376's Process 11 increases the count from 11 to 12.
The Node c0376's Process 12 increases the count from 12 to 13.
The Node c0376's Process 13 increases the count from 13 to 14.
The Node c0376's Process 14 increases the count from 14 to 15.
The Node c0376's Process 15 increases the count from 15 to 16.
The Node c0377's Process 16 increases the count from 16 to 17.
The Node c0377's Process 17 increases the count from 17 to 18.
The Node c0377's Process 18 increases the count from 18 to 19.
The Node c0377's Process 19 increases the count from 19 to 20.
The Node c0377's Process 20 increases the count from 20 to 21.
The Node c0377's Process 21 increases the count from 21 to 22.
The Node c0377's Process 22 increases the count from 22 to 23.
The Node c0377's Process 23 increases the count from 23 to 24.
The Node c0377's Process 24 increases the count from 24 to 25.
The Node c0377's Process 25 increases the count from 25 to 26.
The Node c0377's Process 26 increases the count from 26 to 27.
The Node c0377's Process 27 increases the count from 27 to 28.
The Node c0377's Process 28 increases the count from 28 to 29.
The Node c0377's Process 29 increases the count from 29 to 30.
The Node c0377's Process 30 increases the count from 30 to 31.
The Node c0377's Process 31 increases the count from 31 to 32.
[[li.junce@login-00 q1a]$ █
```

b.)

Please see the details in the code part -> q1/q1b/q1b.c

The result is as follow:

```
[li.junce@login-00 q1b]$ ls
q1b q1b.c q1b.out q1b.script
[li.junce@login-00 q1b]$ more q1b.out
The Node c0205's Process 0 increases the count from 0 to 1.
The Node c0205's Process 1 increases the count from 1 to 2.
The Node c0205's Process 2 increases the count from 2 to 3.
The Node c0205's Process 3 increases the count from 3 to 4.
The Node c0205's Process 4 increases the count from 4 to 5.
The Node c0205's Process 5 increases the count from 5 to 6.
The Node c0205's Process 6 increases the count from 6 to 7.
The Node c0205's Process 7 increases the count from 7 to 8.
The Node c0205's Process 8 increases the count from 8 to 9.
The Node c0205's Process 9 increases the count from 9 to 10.
The Node c0205's Process 10 increases the count from 10 to 11.
The Node c0205's Process 11 increases the count from 11 to 12.
The Node c0205's Process 12 increases the count from 12 to 13.
The Node c0205's Process 13 increases the count from 13 to 14.
The Node c0205's Process 14 increases the count from 14 to 15.
The Node c0205's Process 15 increases the count from 15 to 16.
The Node c0269's Process 16 increases the count from 16 to 17.
The Node c0269's Process 17 increases the count from 17 to 18.
The Node c0269's Process 18 increases the count from 18 to 19.
The Node c0269's Process 19 increases the count from 19 to 20.
The Node c0269's Process 20 increases the count from 20 to 21.
The Node c0269's Process 21 increases the count from 21 to 22.
The Node c0269's Process 22 increases the count from 22 to 23.
The Node c0269's Process 23 increases the count from 23 to 24.
The Node c0269's Process 24 increases the count from 24 to 25.
The Node c0269's Process 25 increases the count from 25 to 26.
The Node c0269's Process 26 increases the count from 26 to 27.
The Node c0269's Process 27 increases the count from 27 to 28.
The Node c0269's Process 28 increases the count from 28 to 29.
The Node c0269's Process 29 increases the count from 29 to 30.
The Node c0269's Process 30 increases the count from 30 to 31.
The Node c0269's Process 31 increases the count from 31 to 32.
The Node c0205's Process 0 decrease the count from 32 to 31.
The Node c0205's Process 1 decrease the count from 31 to 30.
The Node c0205's Process 2 decrease the count from 30 to 29.
The Node c0205's Process 3 decrease the count from 29 to 28.
The Node c0205's Process 4 decrease the count from 28 to 27.
The Node c0205's Process 5 decrease the count from 27 to 26.
The Node c0205's Process 6 decrease the count from 26 to 25.
The Node c0205's Process 7 decrease the count from 25 to 24.
The Node c0205's Process 8 decrease the count from 24 to 23.
The Node c0205's Process 9 decrease the count from 23 to 22.
The Node c0205's Process 10 decrease the count from 22 to 21.
The Node c0205's Process 11 decrease the count from 21 to 20.
The Node c0205's Process 12 decrease the count from 20 to 19.
The Node c0205's Process 13 decrease the count from 19 to 18.
The Node c0205's Process 15 decrease the count from 17 to 16.
The Node c0205's Process 14 decrease the count from 18 to 17.
The Node c0269's Process 16 decrease the count from 16 to 15.
The Node c0269's Process 17 decrease the count from 15 to 14.
The Node c0269's Process 18 decrease the count from 14 to 13.
The Node c0269's Process 19 decrease the count from 13 to 12.
The Node c0269's Process 20 decrease the count from 12 to 11.
The Node c0269's Process 21 decrease the count from 11 to 10.
The Node c0269's Process 22 decrease the count from 10 to 9.
The Node c0269's Process 23 decrease the count from 9 to 8.
The Node c0269's Process 24 decrease the count from 8 to 7.
The Node c0269's Process 25 decrease the count from 7 to 6.
The Node c0269's Process 26 decrease the count from 6 to 5.
The Node c0269's Process 27 decrease the count from 5 to 4.
The Node c0269's Process 28 decrease the count from 4 to 3.
The Node c0269's Process 29 decrease the count from 3 to 2.
The Node c0269's Process 31 decrease the count from 1 to 0.
The Node c0269's Process 30 decrease the count from 2 to 1.
[li.junce@login-00 q1b]$
```

Question 2

a.)

Files "q2a.c" contain the programming used in this question. The scripts "q2a 2nodes.script" and "q2a 4nodes.script" have been used in turn to execute the script with the two different node configurations, 2 and 4, respectively. If we print all of results, it will cost lots of time. Thus, initial results have been obtained from histogramming 1,000 integers between 0 and 200 in 4 bins (using 2 nodes and 2 tasks per node, using q2aTest.c, using the bash script in "q2aTest.script") to verify that the code is operating as intended. And the result is as follow:

```
[li.junce@login-00 q2a]$ more test_q2a_2nodes.out
RESULTS LOCAL[1.000000 - 50.750000) [50.750000 - 100.500000) [100.500000 - 150.250000) [150.250000 - 200.000000)
GLOBAL RESULTS:
[1.000000 - 50.750000) [50.750000 - 100.500000) [100.500000 - 150.250000) [150.250000 - 200.000000)
247 231 269 253
RESULTS LOCAL[1.000000 - 50.750000) [50.750000 - 100.500000) [100.500000 - 150.250000) [150.250000 - 200.000000)
GLOBAL RESULTS:
[1.000000 - 50.750000) [50.750000 - 100.500000) [100.500000 - 150.250000) [150.250000 - 200.000000)
247 231 269 253
RESULTS LOCAL[1.000000 - 50.750000) [50.750000 - 100.500000) [100.500000 - 150.250000) [150.250000 - 200.000000)
GLOBAL RESULTS:
[1.000000 - 50.750000) [50.750000 - 100.500000) [100.500000 - 150.250000) [150.250000 - 200.000000)
247 231 269 253
RESULTS LOCAL[1.000000 - 50.750000) [50.750000 - 100.500000) [100.500000 - 150.250000) [150.250000 - 200.000000)
GLOBAL RESULTS:
[1.000000 - 50.750000) [50.750000 - 100.500000) [100.500000 - 150.250000) [150.250000 - 200.000000)
247 231 269 253
```

Calculation time for 100 processes and 100 bins in the histogram of 2e6 numbers between 1 and 1e6

2Nodes	RunnigTime(s)
1	1.548445
2	1.448345
3	1.543441
4	1.528245
5	1.528349
6	1.445448
7	1.528642
8	1.482447
4Nodes	RunnigTime(s)
1	0.717249
2	0.707238
3	0.725145
4	0.713249
5	0.707248
6	0.717243

4Nodes	RunnigTime(s)
7	0.704248
8	0.718845

A total of 8 experiments, with 2 or 4 nodes in each configuration, have been carried out to assess the code's performance. The below table presents the findings. We see that, on average, the code speeds up by 53.6% in the 4 node configuration compared to the 2 node configuration. This may be due to the fact that communication between nodes is not the bottleneck but rather the computation of the histogram; as a result, using more nodes still results in a faster overall performance.

b.)

The code in "q2a/q2a.c" has been slightly updated for this area of the assignment and saved in a new file called "q2b/q2b.c." Only 20 bins are used to construct the histogram with this version of the code. When there are more processes than bins, control sequences are also included (for example, when there are 4 nodes, there are 3 minimum tasks per node, thus there are 12 processes). If not managed properly, two of these processes—which in this case are idle while the other processes compute the histogram—could cause the application to crash. We don't need to test this code like q2aTest.c because it works fine and is identical to q2a/q2a.c.

Using the files "q2b 2nodes.script" and "q2b 4nodes.script," the identical eight run test has been conducted for each setup. The results are listed below.

Calculation time for 20 processes and 20 bins in the histogram of $2e6$ numbers between 1 and $1e6$

2Nodes	RunnigTime(s)
1	0.162674
2	0.159673
3	0.161649
4	0.161824
5	0.159124
6	0.158674
7	0.162621
8	0.161614
4Nodes	RunnigTime(s)
1	0.217819
2	0.227418
3	0.197919
4	0.209818

4Nodes	RunnigTime(s)
5	0.219817
6	0.216119
7	0.220869
8	0.216817

This time, we see that the performance for both configurations is similar, and in the case of 4 nodes, the performance is significantly lower. In this case, finding a bin only requires each process to iterate through a maximum of 20 bins for each of the 2 million variables. This indicates that each process is under substantially less computational stress and that the communications between nodes are currently the bottleneck. An average 25.3% difference in computing time for the 4-node case reflects this constraint.

c.)

As stated, because there are more processes in part a than in part b, each process in part a must discover the bin for a lesser number of numbers. Finding the bin of each integer, however, requires checking no more than 100 bins, which increases the computing workload compared to the 20 bin situation. Because of this, we can see in section a) that the actual binning at each process is the bottleneck, and additional nodes lead to noticeably higher performance. On the other hand, since the computational load is significantly lower for 20 bins, the bottleneck in part b) is caused by communication between nodes. This leads to improved performance when there are only two nodes.

Question 3

Please see the details of codes in [q3/matmulCUDA.cu](#)

The results have been displayed as a table plot.

Using p100 gpu:

Threads	GPU RunnigTime(ms)	CPU RunnigTime(ms)
32	5.342560	4506.275391
64	0.009152	4479.407227
96	0.008544	4394.415039
128	0.008256	4381.976562
160	0.009248	4380.418457

For p100 gpu:

When I reached 128 threads, like in the experiments I conducted with constant input data size, I found that the performance improvement was not particularly obvious. As a result, I can state that my algorithm's strong scalability reaches its maximum at about 128 threads.

Using k80 gpu:

Threads	GPU RunnigTime(ms)	CPU RunnigTime(ms)
32	21.963680	4185.623535
64	0.013248	4194.656738
96	0.012576	4193.613281
128	0.013184	4193.234375
160	0.012768	4190.564941

For k80 gpu:

Similar to my trials with fixed input data size, I found that the performance boost was not particularly noticeable when I reached 96 threads. I can therefore say that the strong scalability of my method reaches its limit at roughly 96 threads.

Overall, p100 performs more effectively than k80.

Question 4

The "MPI on Millions of Cores" study primarily discussed three parts of the scalability problem, namely the MPI scalability problem, how to increase scalability, and the parallel algorithms used to scale to millions of cores. Large-scale MPI problems exist in 2010, particularly with regard to memory usage. At that time, two concentrations to improve scalability were memory utilization and the effectiveness of all collective functions. The program's fundamental algorithms are the sole way to enable application scalability.

According to the document "A survey of MPI usage in the US exascale computing project" which was first published in 2018, the three primary areas on which exascale computing at the time concentrated were application development, software technology, and hardware technology [1].

According to another research, a novel parallel programming approach that supports massively parallel performance is proposed [2]. It uses coarse- and fine-grained parallelism over inter- and intra-node.

References: [1] Bernholdt, D. E., Boehm, S., Bosilca, G., Gorentla Venkata, M., Grant, R. E., Naughton, T., ... & Vallee, G. R. (2020). A survey of MPI usage in the US exascale computing project. *Concurrency and Computation: Practice and Experience*, 32(3), e4851. [2] Ashraf, M. U., Eassa, F. A., Albeshri, A. A., & Algarni, A. (2018). Toward exascale computing systems: An energy efficient massive parallel computational model. *International Journal of Advanced Computer Science and Applications*, 9(2).